


Pro Coffee Gear – Agent Architecture

How our AI agents are organized: one instance per person, one shared brain per function, everything in git.



• What you're looking at

Every person at Pro Coffee Gear gets **their own private agent instance** – your own box, your own login, nobody else can see into it. On your box you run **profiles**: one for each space you belong to. Everyone carries the LT profile (the shared leadership view) plus the profile for your function (Operations, CS, Sales, etc.). You hop between them with the profile switcher, bottom-right of the app.

What makes it more than "everyone has a chatbot" is the layer in the middle: **Honcho, the shared brain**. Each function has exactly one brain, and every person in that function – and the central function agents – read and write the same memory. Your box is private; your function's knowledge is shared. That's the whole trick.

• How it builds context over time

- ▶ **Day one – seeded, not blank.** Each function's agent starts with standing context loaded from its Notion intake doc: who you are, your vendors, your recurring issues, your guardrails. It knows Pro Coffee Gear from the first message.
- ▶ **Every conversation teaches it.** Honcho remembers what matters from your work sessions – decisions, patterns, preferences – and gets sharper the more you use it. This memory survives restarts, upgrades, and box migrations.
- ▶ **Knowledge crosses functions automatically.** Daily bridges move the relevant signal one-way between brains: what CS learns about broken products flows to Product; what Ops learns about vendor costs flows to Finance. No meetings, no re-explaining.
- ▶ **Leadership sees the whole picture.** The LT profile surfaces the cross-functional digest – material changes and decisions-needed – without anyone having to write status reports.

• How you use it in your function

- ▶ **Treat it like a senior colleague who never forgets.** Ask it questions about your area, draft customer replies, analyze data, write docs – it answers with your function's full accumulated knowledge behind it.
- ▶ **Have it build real things.** Your agent has a terminal, can write code, and can deploy. Dashboards, scripts, internal tools, automations – describe what you need, and iterate with it until it works.
- ▶ **Correct it once, benefit forever.** When it gets something wrong, say so. That correction becomes part of the function's memory – the next person never hits the same mistake.

● How it's visible, shared, and learned from

- ▶ **Everything built lands in git.** Code, dashboards, and tools an agent creates live in versioned repos the whole team can inspect, reuse, and improve. Nothing is trapped inside one person's chat history, and nothing locks us to one vendor.
- ▶ **Function knowledge is inherited.** When someone new joins your function, they plug into the same shared brain on day one – they start where the function is, not from zero.
- ▶ **Wins propagate sideways.** The bridges push one function's learnings to the functions that need them, so a discovery in CS becomes a fix in Product without anyone playing telephone.
- ▶ **Onboarding is self-verifying.** The roster in Notion shows each person's status: Function = what they should have, Active Profiles = what's actually live, and Match / Mismatch tells us instantly if something's off.

● Day one for you

- Accept your Portal invite → your own box
- Run the one onboarding command Wes sends
- Check your roster row: Match + Onboarded ✓
- Open the switcher: LT + your function

● What your agent already knows

- Your function's intake doc (vendors, issues, rules)
- Everything your function's brain has learned
- Cross-functional signal pushed by bridges
- Who you are and how you like to work

● Sharing & learning

- Built things → git repos (inspectable, reusable)

- Function knowledge → shared brain (inherited)
- Learnings → bridges to other functions
- Big picture → LT digest for leadership

Pro Coffee Gear • Agent Architecture • v1.0 • Questions → Wes